

# Process

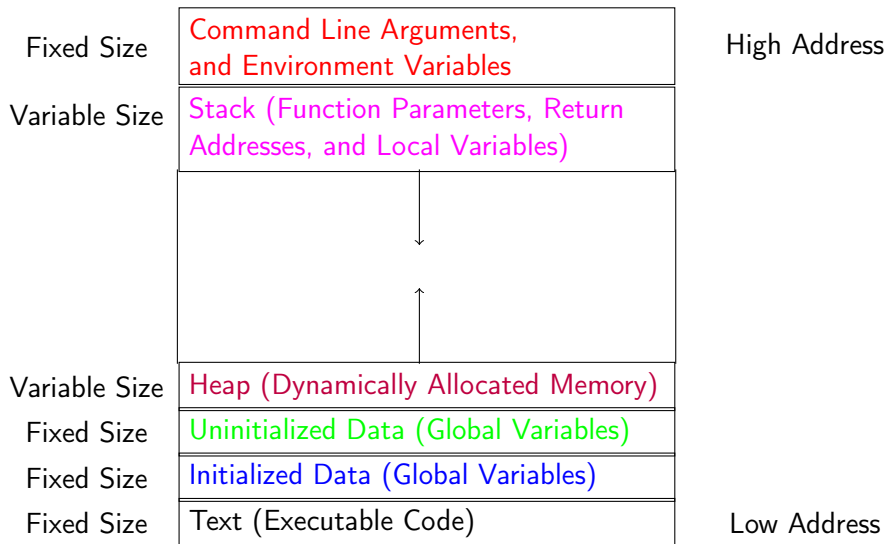
Joy Mukherjee

Indian Institute of Technology Bhubaneswar

# What is a Process?

- A program by itself is not a process.
- A program is a passive entity (an executable file containing a list of instructions stored on disk).
- A **process** is an instance of a program in execution.
- A process is an active entity. The status of the current activity of a process is represented by the value of the program counter and the contents of the CPU registers.
- A program becomes a process when an executable file is loaded into memory.
- Multiple instances of the same program are different processes.
  - Text sections are equivalent.
  - Data, heap, and stack sections vary.

# Memory Layout of a Process



# A C Program

```
#include <stdio.h>
#include <stdlib.h>
int x;                // Uninitialized Data (Global Variables)
int y = 15;           // Initialized Data (Global Variables)
int main(int argc, char *argv[]) { // Command Line Arguments
    int *values;       // Stack (Local Variables)
    int i;
    values = (int *) malloc (sizeof(int) * 5); // Heap
    for(i = 0; i < 5; i++)
        values[i] = i;
    return 0;
}
```

# Process control block (PCB)

- The primary data structure maintained by the kernel that contains information about a process.
- Kernel maintains a list of PCBs for all processes.
- Linux PCB (`task_struct`) has 150+ fields.

# Contents of a PCB

- **Process state:** New / Ready / Running / Waiting / Terminated.
- **Program counter:** The address of the next instruction to be executed for the process.
- **CPU registers:** Accumulators, index registers, stack pointers, and general-purpose registers, plus Program Status Word (PSW).
  - When an interrupt occurs, the program counter, the register contents must be saved to allow the process to be continued correctly afterward, when it is rescheduled to run.

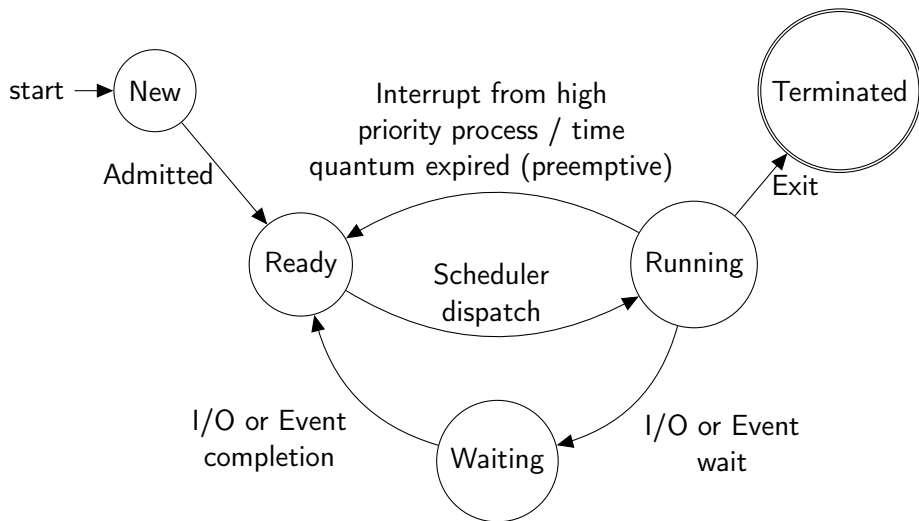
- **Process-scheduling information:** Process priority, pointers to scheduling queues, and any other scheduling parameters.
- **Process-synchronization information:** Sockets, Semaphores, Shared memory regions, Timers, Signal handlers.
- **Memory-management information:** The value of the base and limit registers and the page tables, or the segment tables.
- **Accounting information:** The amount of CPU and real time used, time limits, process id, parent process id, and so on.
- **I/O status information:** The list of I/O devices allocated to the process, a list of open files, and so on.
- Information for each thread.

# States of a Process

- As a process executes, it changes state. The state of a process is defined in part by the current activity of that process.
  - **New**
  - **Ready**
  - **Running**
  - **Waiting / Blocking**
  - **Terminated**
- Only one process can be running on any processor core at any instant.
- Many processes may be in ready and waiting states.



# Process State Diagram



- **New:** The **long term scheduler** brings the process from secondary memory to the ready queue of RAM.
  - The **long term scheduler** controls the degree of multiprogramming (No. of processes in the RAM)
- **Ready:** The process is in the ready queue of RAM, and waiting to be assigned to a CPU based on some CPU scheduling algorithm (FCFS, SJF, SRTF, RR, Priority scheduling).
  - If the ready queue is full, the process is taken to the secondary memory by the **medium term scheduler**.
  - If the ready queue is not full, the process is brought to the ready queue by the **medium term scheduler**.

- **Running:** The **short term scheduler** / **dispatcher** selects the process from the ready queue of RAM and allocates the CPU for execution. In this state, the instructions of the process residing in RAM are being executed.
- **Waiting / Blocking:** The process is waiting for some event to occur such as completion of an I/O or reception of a signal in the waiting queue of RAM.
  - If waiting queue is filled up, the process is taken to the secondary memory by the **medium term scheduler**.
  - If the waiting queue has some space, the process is brought to the waiting queue by the **medium term scheduler**.
- **Terminated:** The process has finished execution. The **long term scheduler** brings the process into the secondary memory from the RAM.

# Main Operations on a Process

- Process creation
- Process scheduling
- Process termination

# Process Creation

- A process (parent) can create another process (child) via `fork()`.
- Each process has a unique ID (an integer).
- PCB is created and initialized
- Initial resources allocated and initialized if needed
- Process added to ready queue (queue of processes ready to run).
- The new process can in turn create other processes, forming a tree of processes.

# Process Creation

- The **systemd** (earlier, **init**) process (which always has a pid of 1) serves as the root parent process for all user processes.
- The **systemd** process is the first user process created when the system boots.
- Once the system has booted, the **systemd** process creates processes which provide additional services such as **logind**, **udev**, and so on.
- **ps -el** lists all active processes in the system.
- **ps tree** displays a tree of all processes in the system.

- **Resource sharing possibilities:**
  - Parent and children share all resources
  - Children share subset of parent's resources
  - Parent and child share no resources
- **Execution possibilities:**
  - Parent and children execute concurrently
  - Parent waits until children terminate
- **Memory address space possibilities:**
  - Address space of child is a duplicate of parent via `fork()`. Child inherits open files, privileges and scheduling attributes from the parent.
  - Child has a new program loaded into it via `execvp()`, `execvp()` etc.

# Process Termination

- Process executes `exit()` / `abort()` and asks the operating system to terminate it.
- Process encounters a fatal error like arithmetic exception etc.
- Parent may terminate execution of children processes (ex. `kill`). Some possible reasons
  - Child has exceeded allocated resources
  - Task assigned to child is no longer required
  - The parent is exiting, and the operating system may not allow a child to continue if its parent terminates.



- Parent may wait for the termination of a child by using the `wait()` system call.
- Parent obtains the exit status of the child in *status*.
- Returns the process identifier *pid* of the terminated child to the parent.

```
pid_t pid;  
int status;  
pid = wait(&status);
```

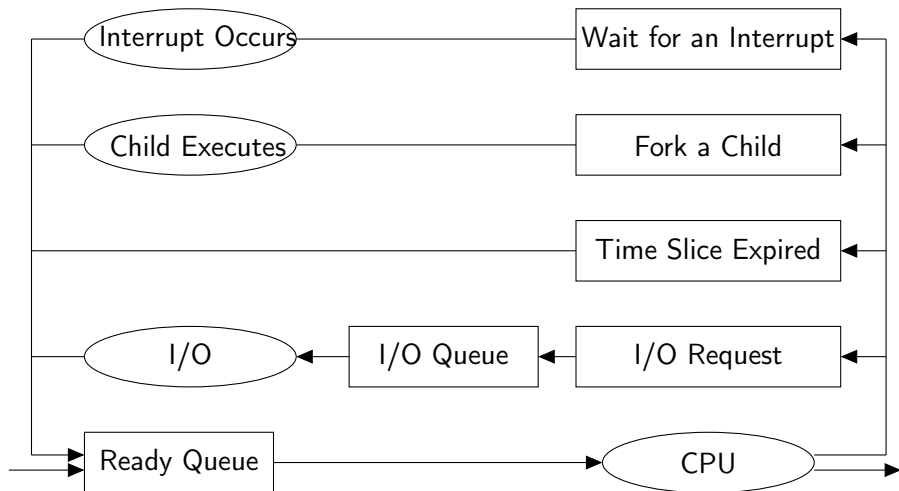
- When a process terminates,
  - All the resources of the process (physical and virtual memory, open files, and I/O buffers) are deallocated by the kernel.
  - However, its entry in the process table must remain there until the parent calls `wait()`, because the process table contains the process's exit status.
  - A process that has terminated, but whose parent has not yet called `wait()`, is known as a **zombie process**.
  - All processes transition to zombie state when they terminate.
  - Once the parent calls `wait()`, the process identifier of the zombie process and its entry in the process table are released.

- If a parent did not invoke `wait()` and terminated,
  - Its child processes are orphans.
  - Linux assigns the **systemd process** (root of the process hierarchy in Linux system) as the new parent to orphan processes.
  - The systemd process periodically invokes `wait()`, thereby allowing the exit status of any orphaned process to be collected and releasing the orphan's process identifier and process-table entry.

# Process Scheduling

- **Ready queue:** queue of PCBs of all processes residing in main memory that are ready to execute
- **Scheduler/Dispatcher:** picks up a process from ready queue according to some CPU Scheduling Policy and assigns it the CPU
- Selected process runs till
  - It needs to wait for some event to occur (ex. a disk read)
  - The CPU scheduling policy dictates that it be stopped
    - CPU time allotted to it expires (timesharing systems)
    - Arrival of a higher priority process
  - When it is ready to run again, it goes back to the ready queue
- Scheduler is invoked again to select the next process from the ready queue

# Queueing Diagram Representation of Process Scheduling



# Scheduling of CPU Scheduler

- The scheduler is scheduled by
  - the code associated with a hardware interrupt, or
  - code associated with a system call

# Context of a Process

- Information that is required to be saved to be able to restart the process later from the same point
- Includes:
  - CPU state – all register contents, PSW
  - Program counter
  - Memory layout of a process
  - Open file information
  - Pending I/O and other event information
- In Linux, each process's context is described by a **task\_struct** structure.

# Context Switch

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process
- Context-switch time is overhead; the system does no useful work while switching
- Context-switch time is highly dependent on hardware support, e.g., the memory speed, the number of registers that must be copied, and the existence of special instructions (such as a single instruction to load or store all registers).
- Context-switch time is typically a few milliseconds.



# Handling Interrupts

- The interrupt descriptor table (IDT) associates each interrupt identifier with an interrupt service routine (ISR).

# Handling Interrupts

- The interrupt descriptor table (IDT) associates each interrupt identifier with an interrupt service routine (ISR).
- Jump to ISR

# Handling Interrupts

- The interrupt descriptor table (IDT) associates each interrupt identifier with an interrupt service routine (ISR).
- Jump to ISR
- ISR first saves the context of the process

# Handling Interrupts

- The interrupt descriptor table (IDT) associates each interrupt identifier with an interrupt service routine (ISR).
- Jump to ISR
- ISR first saves the context of the process
- Execute the ISR

# Handling Interrupts

- The interrupt descriptor table (IDT) associates each interrupt identifier with an interrupt service routine (ISR).
- Jump to ISR
- ISR first saves the context of the process
- Execute the ISR
- Before leaving, ISR should restore the context of the process being executed

# Handling Interrupts

- The interrupt descriptor table (IDT) associates each interrupt identifier with an interrupt service routine (ISR).
- Jump to ISR
- ISR first saves the context of the process
- Execute the ISR
- Before leaving, ISR should restore the context of the process being executed
- Return from ISR restores the PC

# Handling Interrupts

- The interrupt descriptor table (IDT) associates each interrupt identifier with an interrupt service routine (ISR).
- Jump to ISR
- ISR first saves the context of the process
- Execute the ISR
- Before leaving, ISR should restore the context of the process being executed
- Return from ISR restores the PC
- ISR may invoke the dispatcher, which may load the context of a new process, which runs when the interrupt returns instead of the original process interrupted

# Example: Time-sharing Systems

- Each process has a time quantum  $T$  allotted to it



# Example: Time-sharing Systems

- Each process has a time quantum  $T$  allotted to it
- Dispatcher starts process  $P_0$ , loads an external counter (timer) that counts down from  $T$  to 0

# Example: Time-sharing Systems

- Each process has a time quantum  $T$  allotted to it
- Dispatcher starts process  $P_0$ , loads an external counter (timer) that counts down from  $T$  to 0
- When the timer expires, the CPU is interrupted

# Example: Time-sharing Systems

- Each process has a time quantum  $T$  allotted to it
- Dispatcher starts process  $P_0$ , loads an external counter (timer) that counts down from  $T$  to 0
- When the timer expires, the CPU is interrupted
- The ISR invokes the dispatcher

## Example: Time-sharing Systems

- Each process has a time quantum  $T$  allotted to it
- Dispatcher starts process  $P_0$ , loads an external counter (timer) that counts down from  $T$  to 0
- When the timer expires, the CPU is interrupted
- The ISR invokes the dispatcher
- The dispatcher saves the context of  $P_0$

# Example: Time-sharing Systems

- Each process has a time quantum  $T$  allotted to it
- Dispatcher starts process  $P_0$ , loads an external counter (timer) that counts down from  $T$  to 0
- When the timer expires, the CPU is interrupted
- The ISR invokes the dispatcher
- The dispatcher saves the context of  $P_0$
- PCB of  $P_0$  tells where to save

# Example: Time-sharing Systems

- Each process has a time quantum  $T$  allotted to it
- Dispatcher starts process  $P_0$ , loads an external counter (timer) that counts down from  $T$  to 0
- When the timer expires, the CPU is interrupted
- The ISR invokes the dispatcher
- The dispatcher saves the context of  $P_0$
- PCB of  $P_0$  tells where to save
- The dispatcher selects  $P_1$  from ready queue

# Example: Time-sharing Systems

- Each process has a time quantum  $T$  allotted to it
- Dispatcher starts process  $P_0$ , loads an external counter (timer) that counts down from  $T$  to 0
- When the timer expires, the CPU is interrupted
- The ISR invokes the dispatcher
- The dispatcher saves the context of  $P_0$
- PCB of  $P_0$  tells where to save
- The dispatcher selects  $P_1$  from ready queue
- The PCB of  $P_1$  tells where the old state, if any, is saved

# Example: Time-sharing Systems

- Each process has a time quantum  $T$  allotted to it
- Dispatcher starts process  $P_0$ , loads an external counter (timer) that counts down from  $T$  to 0
- When the timer expires, the CPU is interrupted
- The ISR invokes the dispatcher
- The dispatcher saves the context of  $P_0$
- PCB of  $P_0$  tells where to save
- The dispatcher selects  $P_1$  from ready queue
- The PCB of  $P_1$  tells where the old state, if any, is saved
- The dispatcher loads the context of  $P_1$



## Example: Time-sharing Systems

- Each process has a time quantum  $T$  allotted to it
- Dispatcher starts process  $P_0$ , loads an external counter (timer) that counts down from  $T$  to 0
- When the timer expires, the CPU is interrupted
- The ISR invokes the dispatcher
- The dispatcher saves the context of  $P_0$
- PCB of  $P_0$  tells where to save
- The dispatcher selects  $P_1$  from ready queue
- The PCB of  $P_1$  tells where the old state, if any, is saved
- The dispatcher loads the context of  $P_1$
- The dispatcher reloads the counter (timer) with  $T$

## Example: Time-sharing Systems

- Each process has a time quantum  $T$  allotted to it
- Dispatcher starts process  $P_0$ , loads an external counter (timer) that counts down from  $T$  to 0
- When the timer expires, the CPU is interrupted
- The ISR invokes the dispatcher
- The dispatcher saves the context of  $P_0$
- PCB of  $P_0$  tells where to save
- The dispatcher selects  $P_1$  from ready queue
- The PCB of  $P_1$  tells where the old state, if any, is saved
- The dispatcher loads the context of  $P_1$
- The dispatcher reloads the counter (timer) with  $T$
- The ISR returns, restarting  $P_1$  (since  $P_1$ 's PC is now loaded as part of the new context loaded)

## Example: Time-sharing Systems

- Each process has a time quantum  $T$  allotted to it
- Dispatcher starts process  $P_0$ , loads an external counter (timer) that counts down from  $T$  to 0
- When the timer expires, the CPU is interrupted
- The ISR invokes the dispatcher
- The dispatcher saves the context of  $P_0$
- PCB of  $P_0$  tells where to save
- The dispatcher selects  $P_1$  from ready queue
- The PCB of  $P_1$  tells where the old state, if any, is saved
- The dispatcher loads the context of  $P_1$
- The dispatcher reloads the counter (timer) with  $T$
- The ISR returns, restarting  $P_1$  (since  $P_1$ 's PC is now loaded as part of the new context loaded)
- $P_1$  starts running